

Machine Learning (ML)

Chapter 5:

Gradient Descent (GD), Stochastic Gradient Descent (SGD),
and Cross-validation

Saeed Saeedvand, Ph.D.

Outline

In this Chapter:

- ✓ Gradient Descent (GD)
- ✓ Stochastic Gradient Descent (SGD)
- ✓ Stochastic Gradient Descent (Mini-Batch)
- ✓ Cross-Validation approaches
 - K-Fold Cross-Validation
 - Stratified K-Fold Cross-Validation
 - Leave-One-Out Cross-Validation (LOOCV)
 - Time Series Cross-Validation
 - Group K-Fold Cross-Validation

Aim of this chapter:

- ✓ Understanding the Gradient Descent and Stochastic Gradient Descent in practical way as practical optimization algorithms. Understand Cross-Validation approaches solutions and techniques.

What is the Gradient Descent?

- ✓ An optimization algorithm to find the **minimum of a cost function**.
- ✓ Works based on the **slope of the cost function**.
- ✓ We **compute the partial derivatives** of the **cost function** regarding each parameters of the model.

Gradient Descent applications:

- Linear regression:
- Logistic regression (popular classification algorithm)
- Neural networks
- Support vector machines
- Principal component analysis (PCA)
- Clustering
- Recommendation systems
- ...

What is the partial derivative? (Reminder)

- ✓ If we have a function that is constantly **changing** through time (change in input changes output), we use derivative to determine **rate of change**.
- ✓ A **partial derivative** is a derivative, that **shows the change** in only **one chosen variable**.
- ✓ we want to **understand how** the **function changes** with **respect to each variable** while **keeping the other constant**.


In **multivariable functions**, such as $f(x,y)$, for each **independent variable** we calculate **slopes separately** (because the function's behavior can be different in each direction).

$$f(x,y) = x^2y^4, \quad \frac{\partial f(x^2y^4)}{\partial(x)} = 2xy^4 \quad 2(1)(3)^4 = 162$$

With respect to x Slope in x-direction

$$f(x,y) = x^2y^4, \quad \frac{\partial f(x^2y^4)}{\partial(y)} = 4x^2y^3 \quad 4 \times 2^2 \times 3^3 = 432$$

With respect to y Slope of the function in the y-direction at a given point (x, y).



If we **assemble partial derivatives** of all variables **into a matrix** we call it the **Jacobian matrix**.

Gradient Descent (GD)

How to apply the GD?

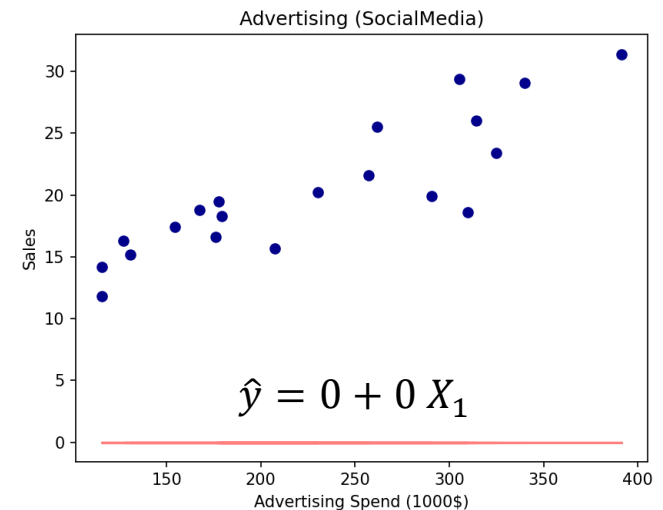
- ✓ In linear regression we needed to estimate the coefficients $\hat{\theta}_0$ and $\hat{\theta}_1$ for the model.

$$\hat{y} = \hat{\theta}_0 + \hat{\theta}_1 x_1$$

Step 1

Initialize the parameters randomly:

- Choose initial values for $\hat{\theta}_0$ and $\hat{\theta}_1$
- For example, $\hat{\theta}_0 = 0$ and $\hat{\theta}_1 = 0$



Gradient Descent (GD)

How to apply the GD?



Step 2 Calculate the cost function:

- ✓ We can Calculate different cost functions: RSS, MSE, ...
 - Here compute the (MSE) between the \hat{y} and y for currenting values of $\hat{\theta}_0$ and $\hat{\theta}_1$.
 - The MSE is given by:

$$MSE = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i]^2$$

Loss or cost function

$$J(\hat{\theta}_0, \hat{\theta}_1) = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i]^2$$



Gradient Descent (GD)

How to apply the GD?

Step 3 Calculate the gradient:

- ✓ We calculate the **partial derivatives** of the cost function regarding $\hat{\theta}_0$ and $\hat{\theta}_1$ independently, (we called it **gradient** witch we assemble slopes into a vector).

Partial derivative of the cost function with respect to $\hat{\theta}_0$, or when it changes (sign is important)

$$\nabla J(\hat{\theta}) = \frac{\partial J}{\partial(\hat{\theta})}$$

Gradient

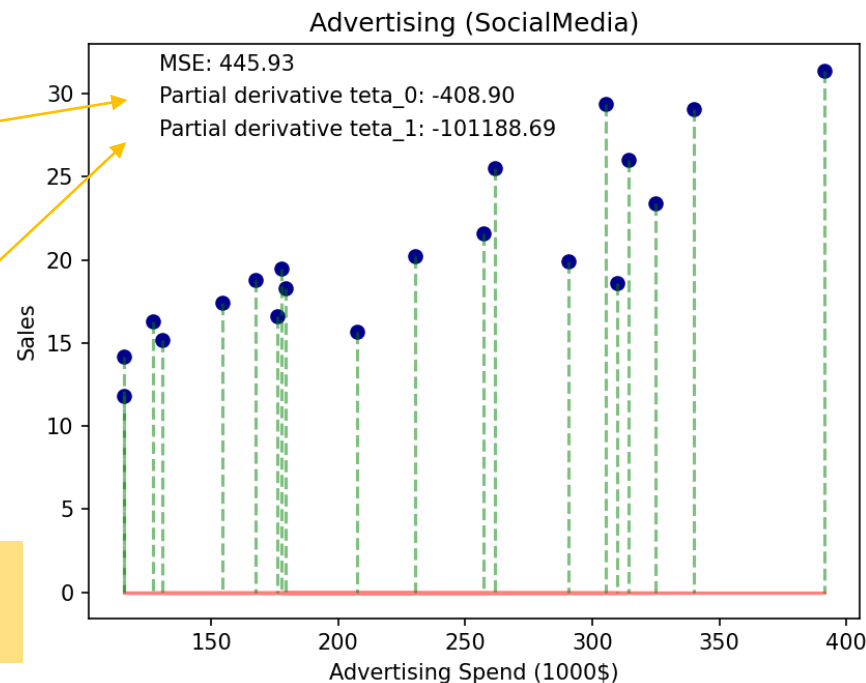
This is simplified result, (since we have learning rate α we can ignore constant -2).

$$\frac{\partial J}{\partial(\hat{\theta}_0)} = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i]$$

$$\frac{\partial J}{\partial(\hat{\theta}_1)} = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i] x_{i_1}$$

$$J(\hat{\theta}_0, \hat{\theta}_1) = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i]^2 \rightarrow J(\hat{\theta}_0, \hat{\theta}_1) = \frac{1}{n} \sum_{i=1}^n [y_i - (\hat{\theta}_0 + \hat{\theta}_1 x_1)]^2$$

$\hat{y} = \hat{\theta}_0 + \hat{\theta}_1 x_1$



Note: these term should be defined based on the objective function and model that we have here MSE).

Gradient Descent (GD)

How to apply the GD?

Step 4 Update the parameters:

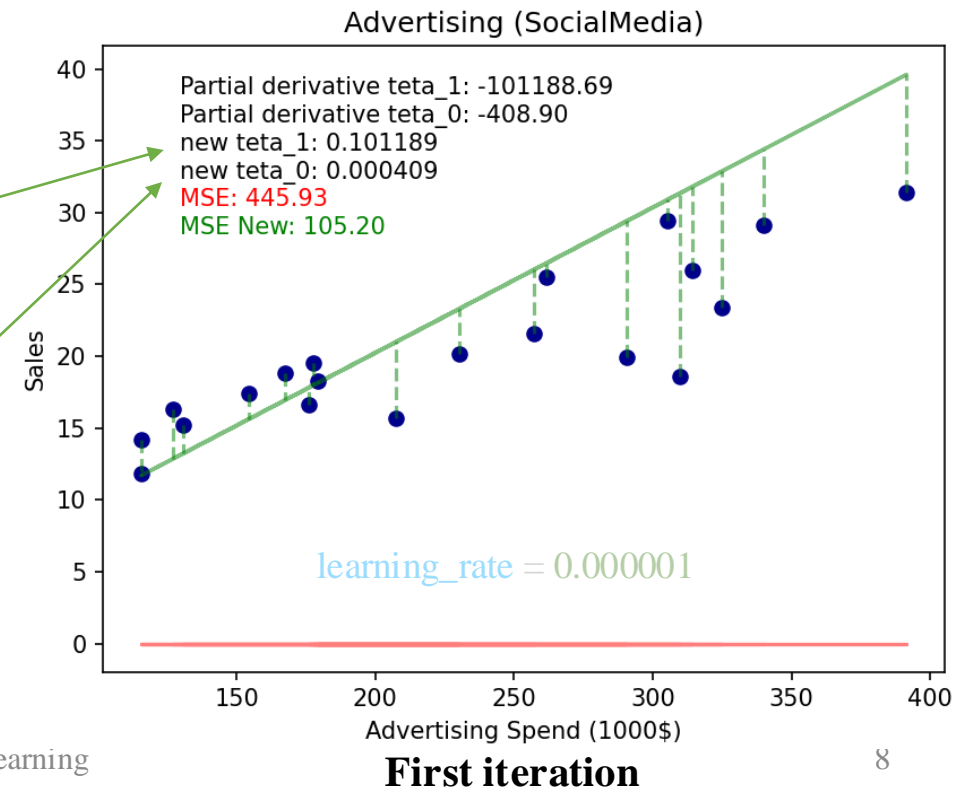

- ✓ We Update the values of $\hat{\theta}_0$ and $\hat{\theta}_1$ using the gradients and a learning rate α :

New $\hat{\theta}_0$ and $\hat{\theta}_1$

$$\hat{\theta}_0 = \hat{\theta}_0 - \alpha \frac{\partial J}{\partial(\hat{\theta}_0)}$$

$$\hat{\theta}_1 = \hat{\theta}_1 - \alpha \frac{\partial J}{\partial(\hat{\theta}_1)}$$

$$\frac{\partial J}{\partial(\hat{\theta}_0)} = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i]$$
$$\frac{\partial J}{\partial(\hat{\theta}_1)} = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i] x_{i,1}$$

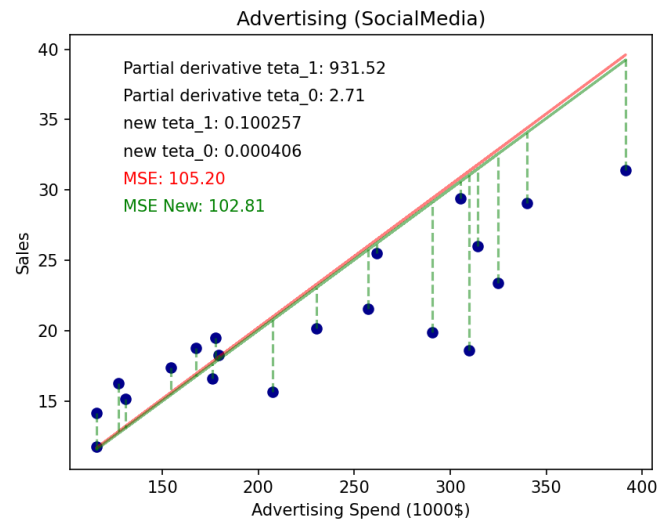


Gradient Descent (GD)


How to apply the GD?

Step 5 Repeat:

- ✓ Run steps 2 to 4 until the cost function gets a minimum or a stopping criterion is met (changes are very small).




Second iteration



Third iteration

...



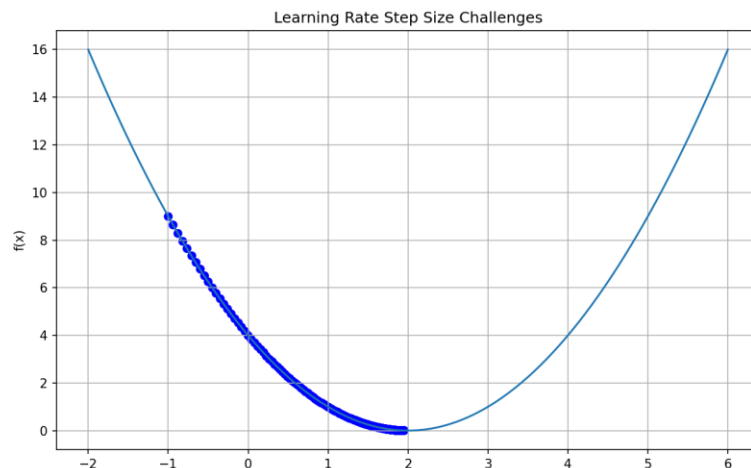
~15th iteration

Gradient Descent (GD)

Learning Rate Challenges

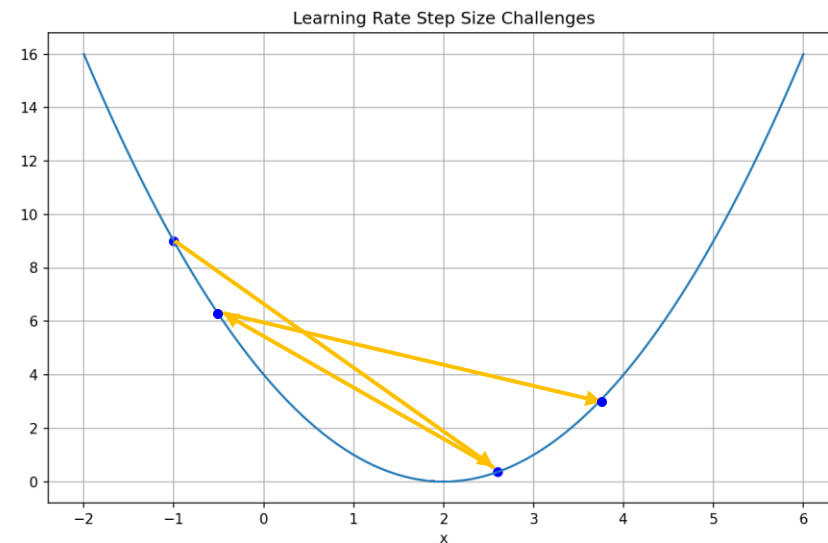
The Small learning rate:

- Converges slowly and can stuck in local minima



Large Learning Rate:

- Overshoot, become unstable and diverge



Stable learning rates: can avoid local minima and converges smoothly

Gradient Descent (GD)


Learning Rate Challenges

✓ For our example if we set big learning rate what happens:

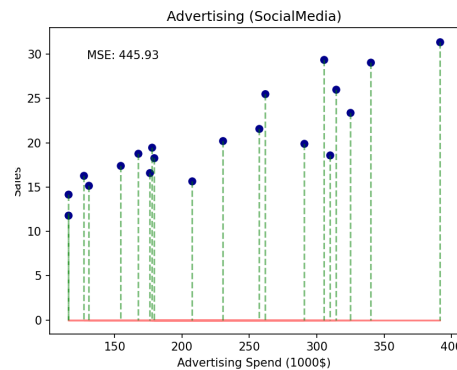
Lets try a Bad learning rate for our example:

`learning_rate = 0.0001`

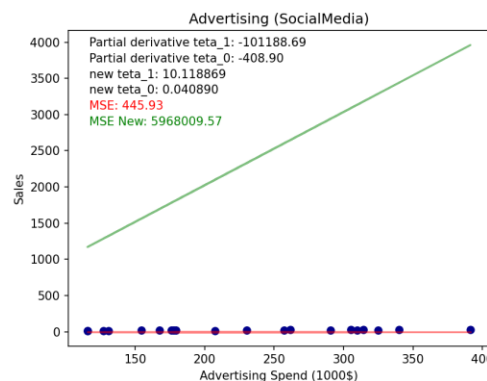
What was Good learning rate?




First iteration



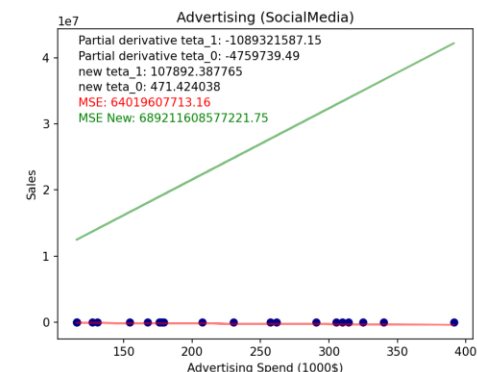
Initial



First iteration



second iteration



Third iteration

...

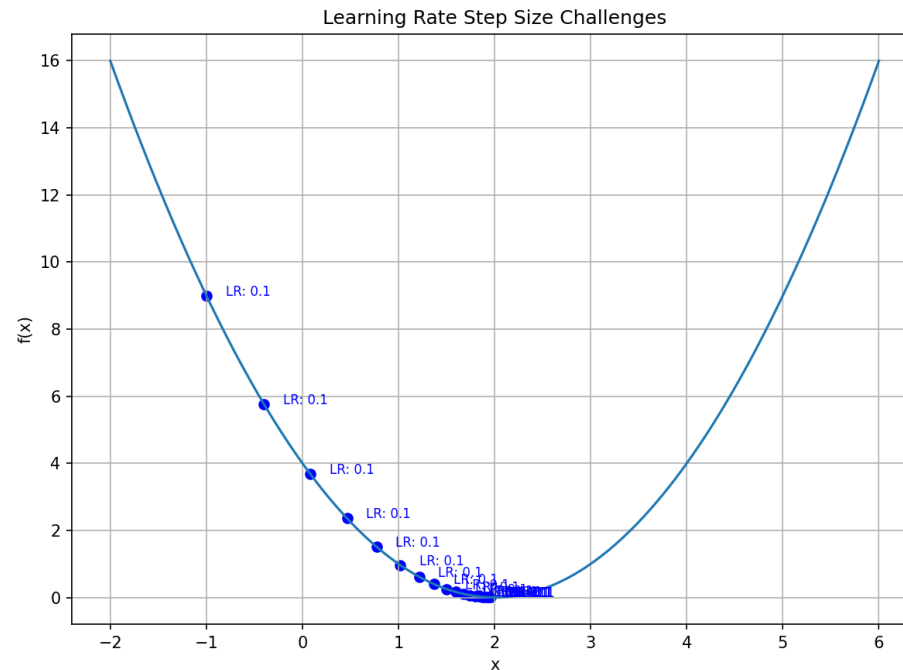
Python Example

Gradient Descent (GD)

Learning Rate Challenges

Proper learning rate:

- Converges slowly enough to find optimal solution.



Gradient Descent (GD)

GD on non-linear regression problem

Step 3
and 4

Calculate the gradient:

✓ All steps same except partial derivatives of the cost function.

$$\hat{y} = \hat{\theta}_0 + \hat{\theta}_1 x_{i-1} + (\hat{\theta}_2 x_{i-1})^2$$

Partial derivatives of the cost function

$$\frac{\partial J}{\partial(\hat{\theta}_0)} = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i]$$

$$\frac{\partial J}{\partial(\hat{\theta}_1)} = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i] x_{i-1}$$

$$\frac{\partial J}{\partial(\hat{\theta}_2)} = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i] (x_{i-1})^2$$



Update of coefficients

$$\hat{\theta}_0 = \hat{\theta}_0 - \alpha \frac{\partial J}{\partial(\hat{\theta}_0)}$$

$$\hat{\theta}_1 = \hat{\theta}_1 - \alpha \frac{\partial J}{\partial(\hat{\theta}_1)}$$

$$\hat{\theta}_2 = \hat{\theta}_2 - \alpha \frac{\partial J}{\partial(\hat{\theta}_2)}$$

Gradient Descent (GD)

Summarize GD Algorithm

Initialize weights (random)

Loop until convergence:

Compute gradient $\frac{\partial J(\hat{\theta})}{\partial \hat{\theta}}$

Update weights, $\hat{\theta} \leftarrow \hat{\theta} - \alpha \frac{\partial J(\hat{\theta})}{\partial \hat{\theta}}$

Return weights

Is there any problem with GD?

Computationally expensive for big datasets (calculate for all points)

$$\frac{\partial J}{\partial (\hat{\theta}_0)} = \frac{1}{n} \sum_{i=1}^n [y_i - \hat{y}_i]$$


$\hat{\theta}$ in this general equation is considered as a vector $(\hat{\theta}_0, \hat{\theta}_1, \dots)$

Stochastic Gradient Descent

Solve the Complexity problem by **SGD!**

Initialize weights (random or zero)

Loop until convergence:




Pick single data Point i

Compute gradient $\frac{\partial J(\hat{\theta})}{\partial \hat{\theta}}$

Update weights, $\hat{\theta} \leftarrow \hat{\theta} - \alpha \frac{\partial J(\hat{\theta})}{\partial (\hat{\theta})}$

Return weights



Problem?

Noisy because we use only one sample!


Stochastic Gradient Descent (Mini-Batch)

SGD Mini-Batch Algorithm

Initialize weights (random or zero)

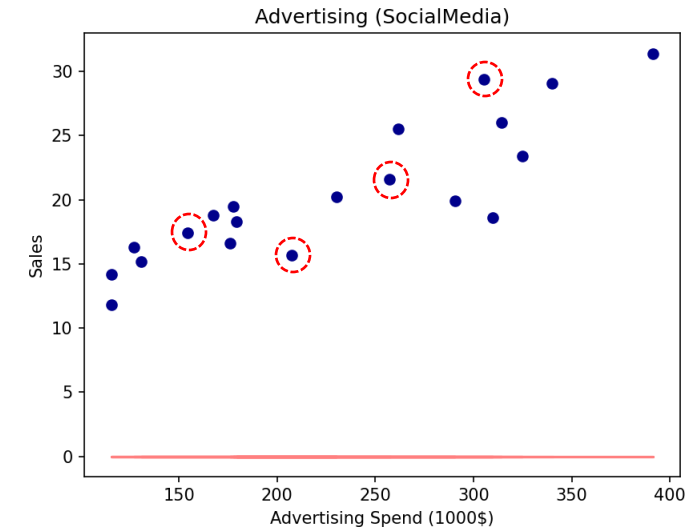
Loop until convergence:

New loop  For each mini-batch of N data points

Compute gradient $\frac{\partial J(\hat{\theta})}{\partial \hat{\theta}}$  **Computationally less expensive**

Update weights, $\hat{\theta} \leftarrow \hat{\theta} - \alpha \frac{\partial J(\hat{\theta})}{\partial \hat{\theta}}$

Return weights



Benefits

- ✓ Smooth convergence.
- ✓ Higher accuracy for estimating Gradient.
- ✓ We can increase learning rate.
- ✓ Parallelized computation.

Stochastic Gradient Descent (Mini-Batch)

Mini-batch Selection

- ✓ As we saw, in SGD we **update the weights** using the **gradients computed** for a small subset (mini-batch) of the dataset.
 1. **Shuffle the training dataset** (to ensure we get more uniform random).
 2. **Chose** or **divide** the dataset into **mini-batches of size**: 16, 32, 64, ..., 512 (the most common ones).

What is the Idea of having multiped baches?

- ✓ **Iterating over the data multiple times** can lead to have **better results**.
- ✓ Hardware and memory **limitations**.

Practice

Run SGD on given example, then change to:
A) quadratic, B) multivariate

Cross-Validation

Definition

- ✓ The **technique** to **assessing the performance** of **machine learning models**.
- ✓ To evaluate **how well** a model **generalizes to unseen or new data**.
- ✓ In cross-validation, we **divide the dataset is into a number of smaller subsets** and we name this **folds**.
- ✓ Generally **we train model** and **test it multiple times with different folds** as **test set (validation set)** and **train set**.
- ✓ We **replot the average performance** of the model **over all iterations** (metrics we already studied like RSS, MSE, p-value, R-squared, ...)

Cross-Validation

Different Approaches (common ones to use):

- ✓ K-Fold Cross-Validation
- ✓ Stratified K-Fold Cross-Validation
- ✓ Leave-One-Out Cross-Validation (LOOCV)
- ✓ Time Series Cross-Validation
- ✓ Group K-Fold Cross-Validation
- ✓ ..

Cross-Validation

k-fold cross-validation

✓ k-fold cross-validation is one of **the most common approach** of Cross-Validation statistical technique.

1. **Divide** the dataset into **k equally sized** folds.
2. **Train** the model on **k-1 folds**.
3. **Test** the model on **remaining one fold**.
4. **Repeat this k times**, so that **each fold being used as a test** set once.
5. **Report the average performance** of the model over the **k iterations**.

Main Advantages

- Can help to fine-tune model hyperparameters.
- Can help to select the best model from a set of candidate models.
- Can help to overcome the risk of overfitting.

k-fold cross-validation

How to choose k?

- ✓ Choosing the value of k for k-fold cross-validation depends on several factors:
 - **Dataset size:** If you have a **small dataset**, choose a **larger value of k** to ensure that each fold has enough data points.
 - **Computational cost:** Increasing the value of k, increases the number of times the model needs to be trained.
 - **Bias-Variance trade-off:** A **smaller k** may lead to **higher variance** in the performance estimate.
 - **Nature of the data:** If the data has a specific structure, you may need to use a **specialized cross-validation technique**.

A common choice for k is 5 or 10, (these values experimentally shown a good balance between computational cost and performance).

Cross-Validation

Challenge of the K-Fold Cross-Validation

- ✓ In K-Fold we divide data into K equally sized folds without considering the distribution of the classes (i.e., the labels).
- ✓ In K-Fold the dataset we may have many samples of a particular class but others only a few (in imbalanced datasets can cause problems).

Cross-Validation

Stratified K-Fold Cross-Validation

- ✓ Similar to K-Fold Cross-Validation but **differ** in how the samples are distributed among the folds.
- ✓ Stratified K-Fold ensures that each fold maintaining the same proportion of samples.
- ✓ So we have **similar distribution** of the data from each class as in the original dataset.

Cross-Validation

Leave-One-Out Cross-Validation (LOOCV)

- ✓ A **special case of K-Fold** Cross-Validation, so that **K equals the number of samples** in the dataset.
- ✓ **At each iteration** **uses a single data point** for validation and the remaining points for training.
- ✓ LOOCV is **computationally expensive**.
- ✓ LOOCV has **higher variance** in performance estimates.
- ✓ LOOCV is **useful in Small datasets**, and **Stable models** (less sensitive to small changes in the training data).

Cross-Validation

Time Series Cross-Validation

- ✓ Specifically designed for time series data, **specially** if the order of the samples is important.
- ✓ In each iteration, we train the model up to a certain time point on the data then we validate the model data after that time point.
- ✓ We **repeat** this process by moving this time window forward.



Cross-Validation

Group K-Fold Cross-Validation

- ✓ When the **dataset contains groups** of related samples:
 1. **Split the data into k groups:** First, **identify the groups** within the dataset (for example one hospital data, or one university data).
 2. **Exclude validation set:** **At each iteration**, **exclude one group** as the validation set, and **train the model**.
 3. **Evaluate models:** **calculate the average performance**.
 4. **Repeat:** Run steps 2-4 this for each of the k groups (iterations).
 5. **Report:** Return **total average performance** metric.

Assignment

Use or extend a sample dataset for our advertising problem, apply one of the proper **Cross-Validation approaches** and train **SGD minibatch to fit the data** (use 3 medias).

Summery

- ✓ We understood the Gradient Descent (GD) optimization approach.
- ✓ We extended GD to Stochastic Gradient Descent (SGD).
- ✓ We improved SGD by Mini-Batch technique.
- ✓ We introduced Cross-Validation approaches including:
 - K-Fold Cross-Validation
 - Stratified K-Fold Cross-Validation
 - Leave-One-Out Cross-Validation (LOOCV)
 - Time Series Cross-Validation
 - Group K-Fold Cross-Validation.