

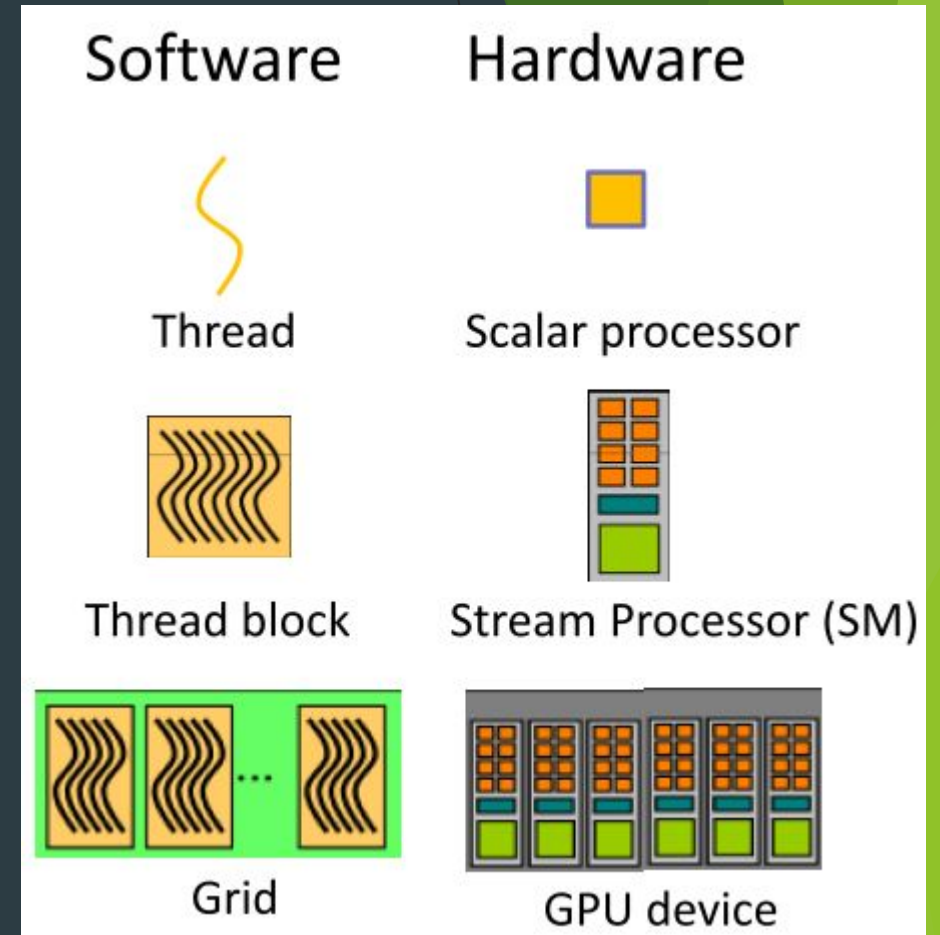
GUDA Introduction II

林政宏

國立台灣師範大學電機系


Execution Model

- ▶ Threads are executed by scalar processor
- ▶ Thread blocks are executed on SM
 - ▶ Several concurrent thread block can reside on one SM
- ▶ A kernel is launched as a grid of thread blocks



Thread Blocks are Executed as Warps

- ▶ Each thread block is mapped to one or more **warps**
 - ▶ When the thread block size is not a multiple of the warp size, unused threads within the last warp are disabled automatically
- ▶ The hardware schedules each warp independently
 - ▶ Warps within a thread block can execute independently




Warps and SIMT

- ▶ A warp is a group of threads within a block that are launched together and (usually) execute together
- ▶ Conceptual Programming Model



- ▶ Conceptual SIMT Execution Model



Warps and SIMT

- ▶ **SIMT = Single Instruction Multiple Threads**
 - ▶ Within CUDA context, refers to issuing a single instruction to the (multiple) threads in a warp.
- ▶ The warp size is currently 32 threads
- ▶ The warp size could change in future GPUs

Filling Warps

- ▶ Prefer thread block sizes that result in mostly full warps
 - ▶ **Bad**: `kernel<<<N, 1>>> (...)`
 - ▶ **Okay**: `kernel<<<N/32, 32>>>(...)`
 - ▶ **Better**: `kernel<<<N/128, 128>>>(...)`
- ▶ Prefer to have **enough threads per block** to provide hardware with many warps to switch between
 - ▶ This is how the GPU hides memory access latency
 - ▶ **128 or 256 threads per block are good choices**

Filling Warps

- ▶ When number of threads is not a multiple of preferred block size, insert bounds test into kernel

```
__global__ void kn( int n, int* x) {  
    int i= threadIdx.x + blockIdx.x * blockDim.x;  
    if (i < n) {  
        //very important  
    }  
}
```

- ▶ Otherwise, threads may access memory outside of arrays
- ▶ Do not launch a second grid to process residual elements 12

```
kernel<<<n/128, 128>>>(...);  
kernel<<<1, n%128>>>(...); // !!! very bad !!!  
kernel<<<(n+127)/128, 128>>>(...); // !!! better!!!
```

Control Flow Divergence

- ▶ Consider the following code

```
__global__ void odd_even(int n, int* x) {  
    int i = threadIdx.x + blockDim.x * blockIdx.x;  
    if ((i & 0x01) == 0) {  
        x[i] = x[i] + 1;  
    }  
    else {  
        x[i] = x[i] + 2;  
    }  
}
```

- ▶ Half the threads in the warp must execute the if clause, the other half the else clause

Performance of Divergent Code

- ▶ Performance decreases with degree of divergence in warps

```
__global__ void dv(int* x)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    switch (i % 32)
    {
        case 0 : x[i] = a(x[i]);
                break;
        case 1 : x[i] = b(x[i]);
                break;
        ...
        case 31: x[i] = v(x[i]);
                break;
    }
}
```

Performance of Divergence

- ▶ Compiler and hardware can detect when all threads in a warp branch in the same direction
 - ▶ For example, all take the if clause, or all take the else clause
 - ▶ The hardware is optimized to handle these cases without loss of performance
- ▶ The compiler can also compile short conditional clauses to use predicates (bits that conditional convert instructions into null ops)
 - ▶ Avoids some branch divergence overheads, and is more efficient
 - ▶ Often acceptable performance with short conditional clauses

Data Address Divergence

- ▶ Concept is similar to control divergence and often conflated
- ▶ Hardware is optimized for accessing contiguous blocks of global memory when performing loads and stores
 - ▶ Global memory blocks are aligned to multiples of 32,64,128 bytes
 - ▶ If requests from a warp span multiple data blocks, multiple data blocks will be fetched from memory, called **memory coalesce**.
 - ▶ Entire block is fetched even if only a single byte is accessed, which can waste bandwidth
- ▶ Hardware handles divergence within **__shared__** memory more efficiently
 - ▶ Designed to support parallel accesses from all threads in warp
 - ▶ Still need to worry about addresses that map to the same bank

Data Address Divergence

- ▶ Hardware may need to issue multiple loads and stores when a warp accesses addresses that are **far apart**
 - ▶ Conceptually similar to executing the load or store multiple times
- ▶ Global memory accesses are most efficient when all load and store addresses generated within a warp are within the same memory block
 - ▶ For example, when addresses of loads and stores have stride 1 within a warp
 - ▶ Common when array index is a linear function of **threadIdx.x**
- ▶ **Consider both address and control divergence when designing algorithms and optimizing code**

Blocking

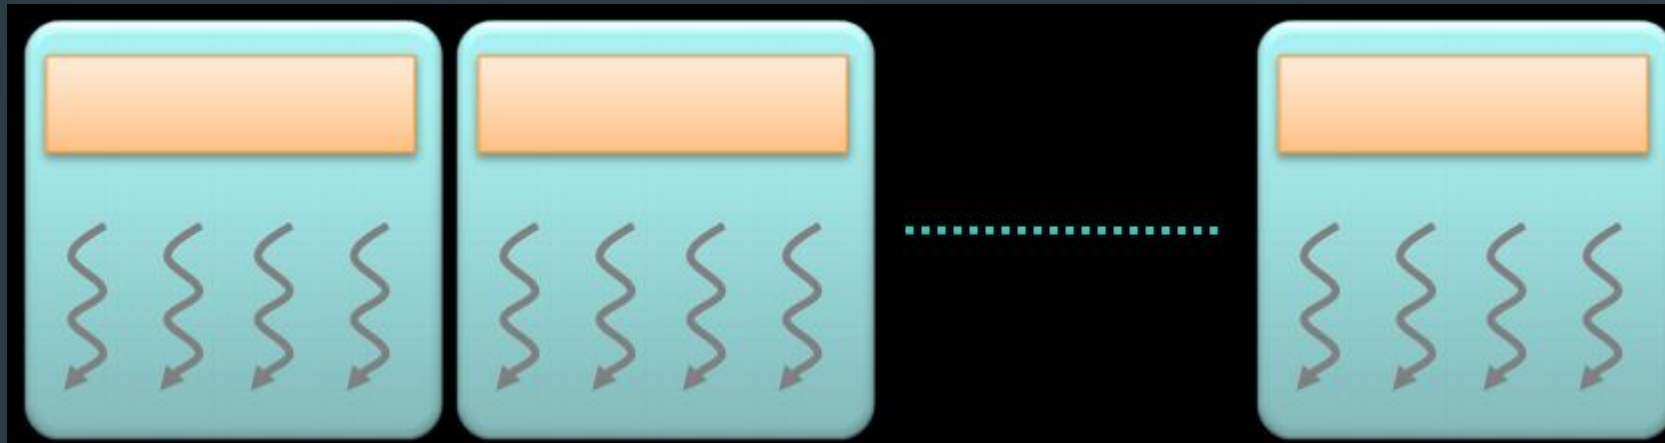
- ▶ Partition data to operate in well-sized blocks
 - ▶ Small enough to be staged in shared memory
 - ▶ Assign each data partition to a thread block
 - ▶ No different from cache blocking!
- ▶ Provides several significant performance benefits
 - ▶ Have **enough thread blocks** to keep processors busy
 - ▶ Working in **shared memory** reduces memory latency dramatically
 - ▶ More likely to have address access patterns that **coalesce** well on load/store to shared memory

Blocking

- ▶ Partition data into subsets that fit into __shared__ memory




- ▶ Process each data subset with one thread block




Blocking

- ▶ Load the subset from global memory to shared memory, using multiple threads to exploit **memory-level parallelism**




Blocking

- ▶ Perform the computation on the subset from shared memory



Blocking

- Copy the result from __shared__ memory back to global memory



Blocking

- ▶ Almost all CUDA kernels are built this way
 - ▶ Blocking may not impact the performance of a particular problem, but one is still forced to think about it
 - ▶ Not all kernels require `__shared__` memory
 - ▶ All kernels do require **registers**

Synchronization

- ▶ Communication
- ▶ Race conditions
- ▶ Synchronizing accesses to shared data

Sharing Data Between Threads

- ▶ Terminology: within a block, threads share data via **shared memory**
- ▶ Extremely fast **on-chip memory**, user-managed
- ▶ Declare using **__shared__**, allocated per block
- ▶ Data is not visible to threads in other blocks

__syncthreads()

- ▶ void __syncthreads();
- ▶ Synchronizes all threads **within** a block
 - ▶ Used to prevent RAW / WAR / WAW hazards
- ▶ All threads must reach the barrier
 - ▶ In conditional code, the condition must be uniform across the block

Global Communication

- ▶ Device threads communicate through shared memory locations
- ▶ Threads in different blocks and different grids
 - ▶ Locations in **global** memory (global variables)
- ▶ Threads in same blocks
 - ▶ Locations in global memory
 - ▶ Locations in shared memory (`__shared__` variables)

Race Conditions

- ▶ Race conditions arise when **2+ threads** attempt to access the same memory location concurrently and at least one access is a write.

```
// race.cu
__global__ void race(int* x)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    *x = i;
}

// main.cpp
int x;
race<<<1,128>>>(d_x);
cudaMemcpy(&x, d_x, sizeof(int), cudaMemcpyDeviceToHost);
```

Race Conditions

- ▶ Programs with **race conditions** may produce unexpected, seemingly arbitrary results
 - ▶ Updates may be missed, and updates may be lost

```
// race.cu
__global__ void race(int* x)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    *x = *x + 1;
}

// main.cpp
int x;
race<<<1,128>>>(d_x);
cudaMemcpy(&x, d_x, sizeof(int), cudaMemcpyDeviceToHost);
```


Synchronization

- ▶ Accesses to shared locations need to be correctly **synchronized** (coordinated) to avoid race conditions
- ▶ In many common shared memory multithreaded programming models, one uses coordination objects such as locks to synchronize accesses to shared data
- ▶ CUDA provides several scalable synchronization mechanisms, such as efficient **barriers** and **atomic** memory operations.
- ▶ In general, always most efficient to design algorithms to avoid synchronization whenever possible.

Synchronization

- Assume thread T1 reads a value defined by thread T0

```
// update.cu
__global__ void update_race(int* x, int* y)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if (i == 0) *x = 1;
    if (i == 1) *y = *x;
}

// main.cpp
update_race<<<1,2>>>(d_x, d_y);
cudaMemcpy(&y, d_y, sizeof(int), cudaMemcpyDeviceToHost);
```

- Program needs to ensure that thread T1 reads location after thread T0 has written location.

Synchronization within Block

- Threads in same block: can use `__syncthreads()` to specify synchronization point that orders accesses

```
// update.cu
__global__ void update(int* x, int* y)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if (i == 0) *x = 1;
    __syncthreads();
    if (i == 1) *y = *x;
}

// main.cpp
update<<<1,2>>>(d_x, d_y);
cudaMemcpy(&y, d_y, sizeof(int), cudaMemcpyDeviceToHost);
```

- Important: all threads within the block must reach the `__syncthreads()` statement

Synchronization between Grids

- Threads in different grids: system ensures writes from kernel happen before reads from subsequent grid launches.

```
// update.cu
__global__ void update_x(int* x, int* y)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if (i == 0) *x = 1;
}

__global__ void update_y(int* x, int* y)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if (i == 1) *y = *x;
}

// main.cpp
update_x<<<1,2>>>(d_x, d_y);
update_y<<<1,2>>>(d_x, d_y);
cudaMemcpy(&y, d_y, sizeof(int), cudaMemcpyDeviceToHost);
```

Synchronization within Grid

- ▶ Often not reasonable to split kernels to synchronize reads and writes from different threads to common locations
 - ▶ Values of `__shared__` variables are lost unless explicitly saved
 - ▶ Kernel launch overhead is non-trivial, and introducing extra launches can degrade performance
- ▶ CUDA provides **atomic** functions (commonly called atomic memory operations) to enforce atomic accesses to shared variables that may be accessed by multiple threads
- ▶ Programmers can synthesize various coordination objects and synchronization schemes using atomic functions.

Introduction to Atomics

- ▶ Atom memory operations (**atomic** functions) are used to solve all kinds of synchronization and coordination problems in parallel computer systems.
- ▶ General concept is to provide a mechanism for a thread to update a memory location such that the update appears to happen atomically (without interruption) with respect to other threads.
- ▶ This ensures that all atomic updates issued concurrently are performed (often in some unspecified order) and that all threads can observe all updates.

Atomic Functions

- Atomic functions perform **read-modify-write** operations on data residing in global and shared memory

```
//example of int atomicAdd(int* addr, int val)
__global__ void update(unsigned int* x)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    int j = atomicAdd(x, 1);    // j = *x; *x = j + i;
}

// main.cpp
int x = 0;
cudaMemcpy(d_x, x, cudaMemcpyHostToDevice);
update<<<1,128>>>);
cudaMemcpy(&x, d_x, cudaMemcpyHostToDevice);
```

- Atomic functions guarantee that only one thread may access a memory location while the operation completes

Atomic Functions

- ▶ The name **atomic** is used because the update is performed atomically: it cannot be interrupted by other atomic updates.
- ▶ The order in which concurrent atomic updates are performed is not defined, and may appear arbitrary. However, none of the atomic updates will be lost.
- ▶ Many different kinds of atomic operations
 - ▶ Add (add), Sub (subtract), Inc (increment), Dec (decrement)
 - ▶ And (bit-wise and), Or (bit-wise or) , Xor (bit-wise exclusive or)
 - ▶ Exch (Exchange)
 - ▶ Min (Minimum), Max (Maximum)
 - ▶ Compare-and-Swap

Histogram Example

```
// Compute histogram of colors in an image
//
// color - pointer to picture color data
// bucket - pointer to histogram buckets, one per color
//
__global__ void histogram(int n, int* color, int* bucket)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if (i < n)
    {
        int c = colors[i];
        atomicAdd(&bucket[c], 1);
    }
}
```

Performance Notes

- ▶ Atomics are **slower** than normal accesses (loads, stores)
- ▶ Performance can degrade when many threads attempt to perform **atomic** operations on a small number of locations

Example: Global Min/Max (Naive)

- ▶ Compute maximum across all threads in a grid
- ▶ One can use a single global maximum value, but it will be VERY slow.

```
__global__ void global_max(int* values, int* global_max)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    int val = values[i];
    atomicMax(global_max, val);
}
```

Example: Global Min/Max (Better)

- Introduce **local maximums** and update global only when new local maximum found.

```
__global__ void global_max(int* values, int* global_max,
                           int *local_max, int num_locals)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    int val = values[i];
    int li = i % num_locals;
    int old_max = atomicMax(&local_max[li], val);
    if (old_max < val)
    {
        atomicMax(global_max, val);
    }
}
```

- Reduces frequency at which threads attempt to update the global maximum, reducing competition access to location.

Lessons from global Min/Max

- ▶ Many updates to a **single value** causes serial bottleneck
- ▶ One can create a hierarchy of values to introduce more parallelism and locality into algorithm
- ▶ 但是, 性能仍然可能很慢, 因此請謹慎使用