

Accelerating String Matching on Graphic Processing Units

Cheng-Hung Lin
National Taiwan Normal University
Taipei, Taiwan

Outline

- Introduction
- Parallel Failureless Aho-Corasick Algorithm
- Memory-Efficient Memory Architecture
- M-DFA (Multithreaded DFA) for Regex Matching

String Matching

- String matching engine plays an important role in many applications, such as network intrusion detection systems, spam filters, and bioinformatics.
- String matching is used to find a place where one or several strings (also called patterns) are found within a larger string or text.

Patterns

“TACT”

“CTO”

“TOE”


Input string

T A C T O E

↑ ↑ ↑
T A C T
C T O
T O E

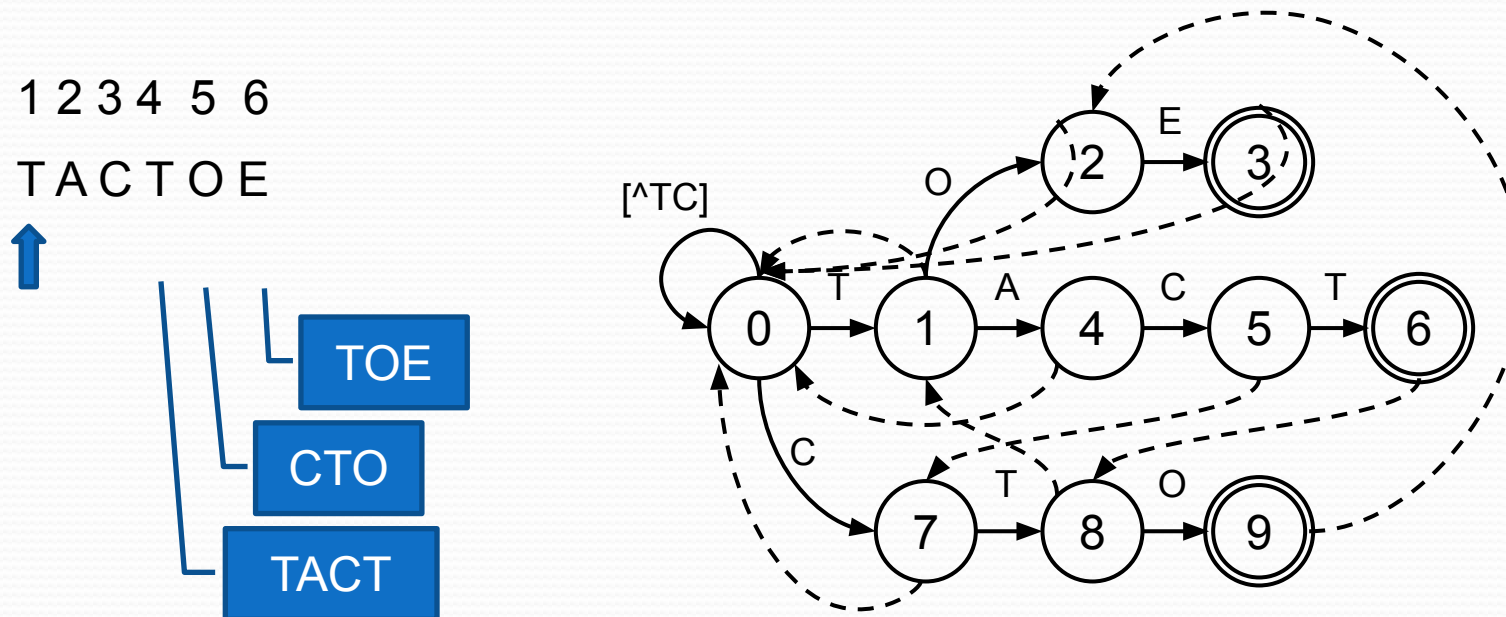
Aho-Corasick Algorithm

- Aho-Corasick algorithm has been widely used for string matching due to its advantage of matching multiple string patterns in a single pass
- Aho-Corasick algorithm compiles multiple string patterns into a state machine




Aho-Corasick Algorithm (cont.)

- String matching is performed by traversing the Aho-Corasick (AC) state machine
- Failure transitions are used to backtrack the state machine to recognize patterns in different locations.




Naïve Data Parallel Approach

- Partition an input stream into multiple segments and assign each segment a thread to traverse AC state machine
- Boundary detection problem
 - Pattern occurs in the boundary of adjacent segments.
 - Duration time of threads = segment size + longest pattern length - 1



Parallel Failureless Aho-Corasick Algorithm


- Parallel Failureless Aho-Corasick (PFAC) algorithm on graphic processing units
 - Allocate each byte of input an individual thread to traverse a state machine





- Reference:
 - C.-H. Lin *et al.*, "Accelerating String Matching Using Multi-Threaded Algorithm on GPU," in *GLOBECOM 2010*.
 - C.-H. Lin *et al.*, "Accelerating Pattern Matching Using a Novel Parallel Algorithm on GPUs", *IEEE Transactions on Computers*

Failureless-AC State Machine


- Remove all failure transitions as well as the self-loop transitions backing to the initial state
 - Minimum number of valid transitions
 - Thread is terminated when no valid transitions




Mechanism of PFAC



Thread# n



Thread# $n+2$



Thread# $n+3$

Experimental Environment


- Intel Core™ i7-950
 - Quad cores
 - 12GB DDR3 memory
- Nvidia® GeForce® GTX580
 - 512 cores
 - 1536MB GDDR5 memory
- Patterns: String pattern extracted from Snort V2.8, containing 27754 states, 1998 final states (patterns)
- Input: extracted from DEFCON

Implementations

- AC_{CPU} : implementation of the AC algorithm on the Core™ i7 using a single thread and optimized by GCC 4.4.3 using the compiler flags “-O2 -msse4”.
- $DPAC_{OMP}$: implementation of the DPAC algorithm on Intel Core™ i7 CPU with OpenMP and optimized by GCC 4.4.3 using the compiler flags “-O2 -msse4”.
- $PFAC_{OMP}$: implementation of the PFAC algorithm on Intel Core™ i7 CPU with the OpenMP and optimized by GCC 4.4.3 using the compiler flags “-O2 -msse4”.
- $PFAC_{GPU}$: implementation of the PFAC algorithm on NVIDIA GPUs.

Performance Evaluation

$$\text{Raw data throughput} = \frac{\text{input_size}}{t_{GPU}}$$



PFAC Library

- **PFAC** is an open source library for multiple string matching performed on **Nvidia GPUs**.
 - PFAC runs on Nvidia GPUs that support **CUDA**, including NVIDIA 1.1, 1.2, 1.3, 2.0 and 2.1 architectures.
 - Supporting OS includes ubuntu, Fedora and MAC OS.
- Released at Google code project
 - <http://code.google.com/p/pfac/>
 - Provides C-style API
 - Users don't need to have background of GPU programming



pfac

PFAC is an open library for exact string matching performed on NVIDIA GPUs

Search projects

[Project Home](#) [Downloads](#) [Wiki](#) [Issues](#) [Source](#) [Administer](#)

[Summary](#) [People](#)

Tip: Project owners, see our [Getting Started](#) guide for steps to configure your project. [X](#)

Project Information

Recommend this on Google+

★ Starred by 8 users
[Project feeds](#)

Code license
[Apache License 2.0](#)

Labels
[Academic](#), [Cuda](#),
[GPUcomputing](#),
[patternmatching](#),
[stringmatching](#),
[parallelcomputing](#),
[dataparallelalgorithm](#)

Members
[brucelinco](#), [LungShengChien](#),
[lgen7604](#), [shihchieh.chang](#)

Your role
[Owner](#)

Links

External links
[CUDA forum](#)
[GPGPU.org](#)
[CUDA Training](#)

Groups
[pfac Forum](#)

What is PFAC?

PFAC is an open library for exact string matching performed on **GPUs**. PFAC runs on processors that support **CUDA**, including NVIDIA 1.1, 1.2, 1.3, 2.0 and 2.1 architectures. Supporting OS includes ubuntu, Fedora and MAC OS.

PFAC library provides C-style API and users need not have background on GPU computing or parallel computing. PFAC has APIs hiding CUDA stuff.

News

- PFAC [r1.0](#) updated 2011/02/23
- PFAC [r1.0](#) released 2011/02/21
- PFAC [r1.1](#) released 2011/04/27
- PFAC [r1.2](#) released 2011/04/29

Simple Example

Example 1: Using PFAC_matchFromHost function

The file "example_pattern" in the directory "../test/pattern/" contains 4 patterns.

```
AB
ABG
BEDE
ED
```

The file "example_input" in the directory "../test/data/" contains a string.

```
ABEDEDABG
```


Five Steps to Use PFAC for String Matching

The following example shows the basic steps to use PFAC library for string matching.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <assert.h>
#include <PFAC.h>

int main(int argc, char **argv)
{
    char dumpTableFile[] = "table.txt" ;
    char inputFile[] = "../test/data/example_input" ;
    char patternFile[] = "../test/pattern/example_pattern" ;
    PFAC_handle_t handle ;
    PFAC_status_t PFAC_status ;
    int input_size ;
    char *h_inputString = NULL ;
    int *h_matched_result = NULL ;

    // step 1: create PFAC handle
    PFAC_status = PFAC_create( &handle ) ;
    assert( PFAC_STATUS_SUCCESS == PFAC_status );

    // step 2: read patterns and dump transition table
    PFAC_status = PFAC_readPatternFromFile( handle, patternFile ) ;
    if ( PFAC_STATUS_SUCCESS != PFAC_status ){
        printf("Error: fails to read pattern from file, %s\n", PFAC_getErrorString(PFAC_status) );
        exit(1) ;
    }

    // dump transition table
    FILE *table_fp = fopen( dumpTableFile, "w" ) ;
    assert( NULL != table_fp ) ;
    PFAC_status = PFAC_dumpTransitionTable( handle, table_fp ) ;
    fclose( table_fp ) ;
    if ( PFAC_STATUS_SUCCESS != PFAC_status ){
        printf("Error: fails to dump transition table, %s\n", PFAC_getErrorString(PFAC_status) );
        exit(1) ;
    }
}
```

```

//step 3: prepare input stream
FILE* fpin = fopen( inputFile, "rb");
assert ( NULL != fpin );

// obtain file size
fseek (fpin , 0 , SEEK_END);
input_size = ftell (fpin);
rewind (fpin);

// allocate memory to contain the whole file
h_inputString = (char *) malloc (sizeof(char)*input_size);
assert( NULL != h_inputString );

h_matched_result = (int *) malloc (sizeof(int)*input_size);
assert( NULL != h_matched_result );
memset( h_matched_result, 0, sizeof(int)*input_size );

// copy the file into the buffer
input_size = fread (h_inputString, 1, input_size, fpin);
fclose(fpin);

// step 4: run PFAC on GPU
PFAC_status = PFAC_matchFromHost( handle, h_inputString, input_size, h_matched_result );
if ( PFAC_STATUS_SUCCESS != PFAC_status ){
    printf("Error: fails to PFAC_matchFromHost, %s\n", PFAC_getErrorString(PFAC_status) );
    exit(1) ;
}

// step 5: output matched result
for (int i = 0; i < input_size; i++) {
    if (h_matched_result[i] != 0) {
        printf("At position %4d, match pattern %d\n", i, h_matched_result[i]);
    }
}

```

The screen shows the following matched results.

```

At position    0, match pattern 1
At position    1, match pattern 3
At position    2, match pattern 4
At position    4, match pattern 4
At position    6, match pattern 2



```


Outline

- Introduction
- Parallel Failureless Aho-Corasick Algorithm
- Memory-Efficient Memory Architecture
- M-DFA (Multithreaded DFA) for Regex Matching

Memory Issue of PFAC

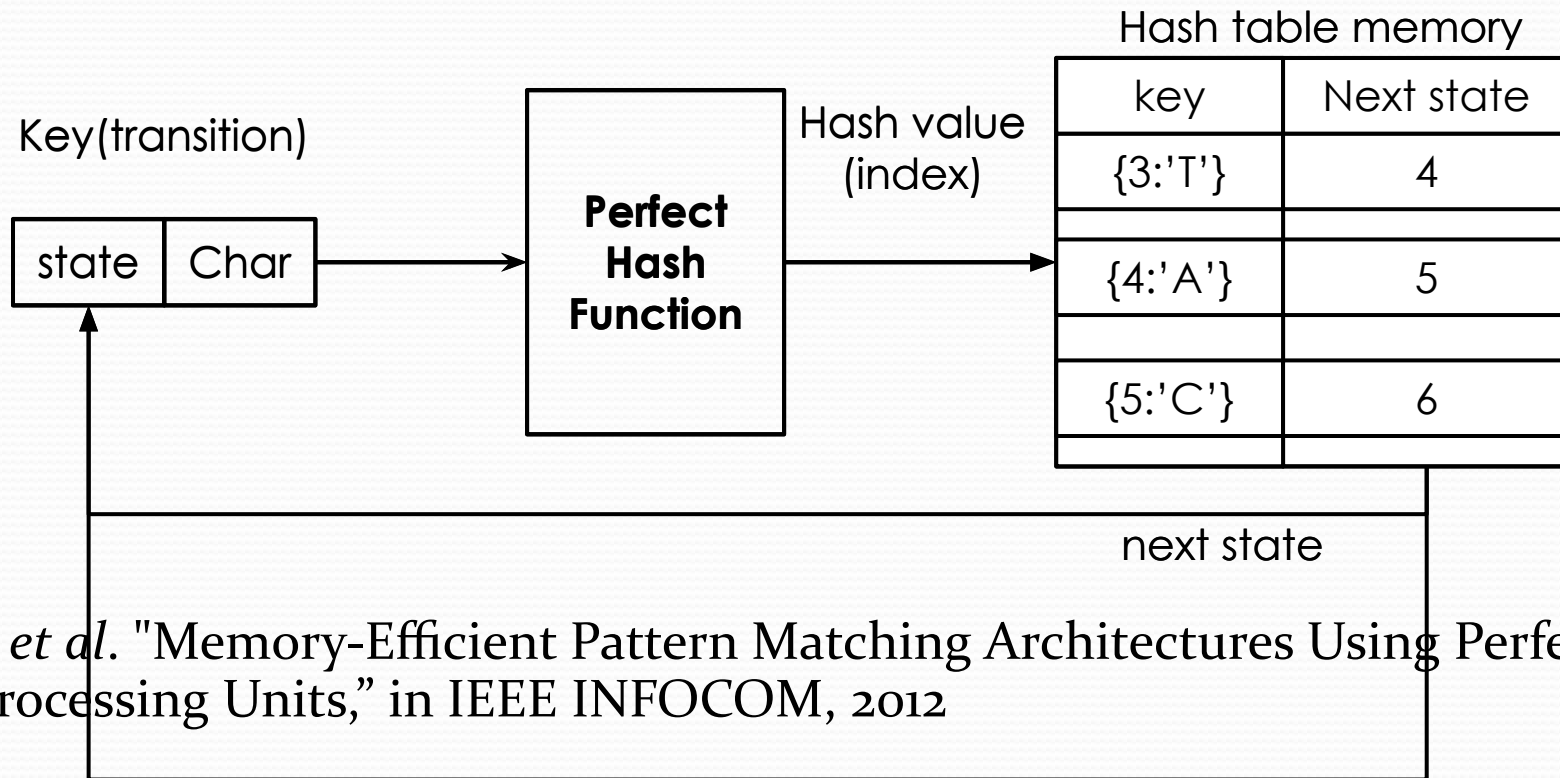

- The two-dimensional memory is sparse.
 - Each row (state) needs 1K (256 x 4)bytes
 - A state machine with 1M states needs 1G bytes
 - 99% of memory is wasted



- Design a **compact storage mechanism** for storing PFAC state transition table is essential for GPU implementation.

Perfect Hashing Memory Architecture

- Use a perfect hash function to store only valid transitions of a PFAC state machine in a hash table



- Reference

- C.-H. Lin, *et al.* "Memory-Efficient Pattern Matching Architectures Using Perfect Hashing on Graphic Processing Units," in IEEE INFOCOM, 2012

Hardware-friendly Perfect Hash Function

- Slide-Left-then-Right First-Fit (SLRFF) algorithm
- Steps to create PHF table
 1. Start with a two-dimensional table of width w and place each key k at location (row, col) , where $row = k / w$, $col = k \bmod w$.
 2. Rows are prioritized by the number of keys in it and move rows in order of priority as following steps.
 - a) First, slide the row left to let the first key in the row be aligned at the first column.
 - b) Then, slide the row right until each column has only one key and record the offset in an array.
 3. Collapse the two-dimensional key table to a linear array.
- Reference
 - R. E. Tarjan and A. C.-C. Yao, "Storing a sparse table," Commun. ACM, vol. 22, pp. 606-611, 1979.

Step 1 of Creating PHF

- Key set, $S = \{2, 4, 10, 11, 13, 14, 17, 20, 21, 25, 27\}$
- Start with a two-dimensional table of width w and place each key k at location $(row, column)$, where $row = k / w$, $column = k \bmod w$.

Key table ($w = 8$)

		2		4			
		10	11		13	14	
	17			20	21		
	25		27				

Step 2 of Creating PHF

- Rows are prioritized by the number of keys in it
- According to the order of priority
 - a) Slide each row left to let the first key be aligned at the first column.
 - b) Slide each row right until each column has only one key and record the offset in the array RT.

3	RT[0] = 9			2		4			
1	RT[1] = 0			10	11		13	14	
2	RT[2] = 0		17			20	21		
4	RT[3] = 6		25		27				

Step 3 of Creating PHF

- Collapse the two-dimensional table to a linear table HK.

RT[0] = 5							2		4			
RT[1] = -2	10	11		13	14							
RT[2] = 1			17			20	21					
RT[3] = 6									25		27	

HK :	10	11	17	13	14	20	21	2	25	4	27
------	----	----	----	----	----	----	----	---	----	---	----

Computation of Hash Value

- $row = k / w$;
- $col = k \bmod w$;
- $index = RT[row] + col$;
- If $HK[index] == k$
 k is a valid key ;
else
 k is an invalid key ;

For example:

Given $k = 14$

$row = 14 / 8 = 1$

$col = 14 \bmod 8 = 6$

$index = RT[1] + 6 = -2 + 6 = 4$

$HK[4] = 14$

14 is a valid key

index : 0 1 2 3 4 5 6 7 8 9 10

HK :	10	11	17	13	14	20	21	2	25	4	27
------	----	----	----	----	----	----	----	---	----	---	----

RT[0] = 5
RT[1] = -2
RT[2] = 1
RT[3] = 6

Computation of Hash Value

- $row = k / w$;
- $col = k \bmod w$;
- $index = RT[row] + col$;
- If $HK[index] == k$
 k is a valid key ;
else
 k is an invalid key ;

For example:

Given $k = 19$

$row = 19 / 8 = 2$

$col = 19 \bmod 8 = 3$

$index = RT[2] + 3 = 1 + 3 = 4$

$HK[4] = 14$

19 is an invalid key

index : 0 1 2 3 4 5 6 7 8 9 10

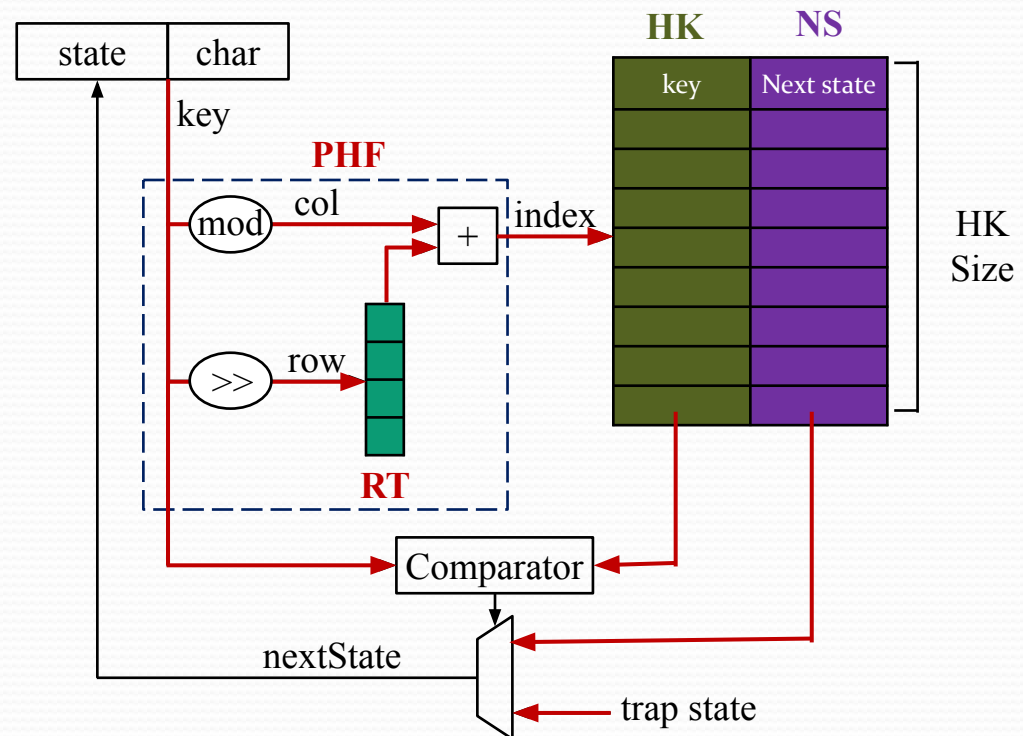
HK :	10	11	17	13	14	20	21	2	25	4	27
------	----	----	----	----	----	----	----	---	----	---	----

RT[0] = 5
RT[1] = -2
RT[2] = 1
RT[3] = 6

Perfect Hashing Memory Architecture

Algorithm

- $row = k / w$;
- $col = k \bmod w$;
- $index = RT[row] + col$;
- If $HK[index] == k$
 k is a valid key ;
else
 k is an invalid key ;
- If k is a valid key
 $nextState = NS[index]$;
else
 $nextState = \text{trap state}$;



Experimental Environment

- Intel Core™ i7-950
 - Quad cores
 - 12GB DDR3 memory
- Nvidia® GeForce® GTX580
 - 512 cores
 - 1536MB GDDR5 memory
- Patterns: String pattern extracted from Snort V2.8, containing 126,776 states, 10,076 final states (patterns)
- Input: 256MB packets extracted from DEFCON

Experimental Results

	# of Rules	# of char	Memory (Bytes)	$\frac{mem}{char}$	Throughput (Gbps)
PHM	10K	187K	620KB	3.39B	100.76
	2K	41K	137KB	3.34B	135.93
PFAC	2K	41K	27.1MB	677B	146.63
B-FSM	39.5K	25.2K	188KB	7.4B	2
CDFA	1,785	29.0K	129KB~ 256KB	4.45B~ 8.2B	11.7
Bitmap Compression	1.5K	18.2K	2.8MB	154B	7.6
Path Compression	1.5K	18.2K	1.1MB	60B	7.6

*PHM : Perfect Hashing Memory Architecture

Conclusions

- The PFAC algorithm is adaptive to be implemented on GPUs and multicore CPUs.
- The perfect hash algorithm significantly reduces the memory for storing state transition table with little penalty on performance.

M-DFA (Multithreaded DFA): An Algorithm for Reduction of State Transitions and Acceleration of REGEXP Matching

Cheng-Hung Lin
National Taiwan Normal
University
Taipei, Taiwan
brucelin@ntnu.edu.tw


Jyh-Charn Liu
Dept. of Computer Science &
Engineering
Texas A&M University
liu@cse.tamu.edu

Regular Expression Matching

- Regular Expression (REGEXP) matching is typically implemented as a **nondeterministic finite automaton (NFA)** or its equivalent **deterministic finite automata (DFA)**.
- NFA has smaller sizes in terms of memory space, but it may take multiple cycles to match an input symbol when multiple states become *active* concurrently.
- An NFA can be mapped to its DFA equivalence by mapping concurrent active states in NFA to one single active state in the DFA.
 - State explosion arises during integration of multiple regexps into a DFA


NFA and DFA

● NFA of “[abc]x” and “c[xy]”



state	Match Patterns
4	“[abc]x”
5	“c[xy]”

● DFA




state	Match Patterns
4	“[abc]x”
5	“c[xy]”
6	“[abc]x” and “c[xy]”

M-DFA (multithreaded DFA)

- Multi-thread based regular expression (regexp) matching algorithm for parallel computer architectures such as multi-core processors and graphic processing units (GPU)
 - Each input symbol is treated as the first symbol of a possibly matched substring to the regexp
 - A DFA free of backtracking transitions is assigned to a thread, which terminates either when it reaches the final state or any mismatch occurs,
 - Multiple threads concurrently read in each input symbol for their independent matching.


Multi-threaded M-DFA Model

- DFA_{LF}



state	Match Patterns
4	"[abc]x"
5	"c[xy]"
6	"[abc]x" and "c[xy]"

- Multi-threaded execution of multiple DFA_{LF}



Construction of DFA_{LF}

- Three steps:

1. Convert regexps into an NFA by Thompson's algorithm
2. Convert the NFA into an equivalent DFA by Rabin-Scott powerset construction algorithm
3. Remove all backtracking transitions of the DFA

Advantages of M-DFA

- M-DFA resembles the behavior of an NFA at the string matching level, yet each thread is a DFA.
- M-DFA is guaranteed to eliminate backtracking transitions.
 - Significant memory reduction

Experimental Results

- M-DFA on Nvidia® GeForce® GTX480 GPU
- RE2 [5] on Intel Core™ i7-950 CPU
- Benchmarks: 16 regular expression patterns and 2 exact string patterns from a public benchmark
- 35 times faster than RE2
- 44% of state reduction and 99.8% of transition reduction

	RE2 library			M-DFA _{GPU}						
Input size (bytes)	# of states	# of state transition	elapsed time(ms)	# of states	state reduction	# of state transition	Transition reduction	GPU elapsed time (ms)	Total elapsed time (ms)	speedup
100K	162	41,472	14.08	91	44%	90	99.8%	0.05	0.35	40.23
32M			2556					5.59	69.24	36.92
64M			4980					11.16	139.32	35.74

Thanks for your attention!

Q&A